

Capstone Project - III

AI



**ML Supervised
Classification**

**Bank Marketing
Effectiveness
Prediction**

Presented by-

Prayag Raj Dubey

CONTENT

AI

Objective

01
STEP

**Data
Summary**

02
STEP

**Exploratory
Data Analysis**

03
STEP

**Feature
Engineering**

04
STEP



05
STEP

**Data
Preparation**

06
STEP

**Handling
Imbalance**

07
STEP

**Classification
Models**

08
STEP

Conclusions

Data Description and Objective

The data is related with direct marketing campaigns (phone calls) of a Portuguese banking institution.

The marketing campaigns were based on phone calls. Often, more than one contact to the same client was required, in order to access if the product (bank term deposit) would be ('yes') or not ('no') subscribed.

The classification goal is to predict if the client will subscribe a term deposit (variable y).



Dataset Attributes and their Description

AI

Bank Client data:

- **age** (numeric)
- **job** : type of job (categorical)
- **marital** : marital status (categorical)
- **education** (categorical)
- **default**: has credit in default? (categorical)
- **housing**: has housing loan? (categorical)
- **loan**: has personal loan? (categorical)

Related with the last contact of the current campaign:

- **contact**: contact communication type (categorical)
- **month**: last contact month of year (categorical)
- **Day of week**: last contact day of the week (categorical)
- **duration**: last contact duration, in seconds (numeric)



Other attributes:

- **campaign**: number of contacts performed during this campaign and for this client (numeric, includes last contact)
- **pdays**: number of days that passed by after the client was last contacted from a previous campaign (numeric)
- **previous**: number of contacts performed before this campaign and for this client (numeric)
- **poutcome**: outcome of the previous marketing campaign (categorical)

Output variable (desired target):

- **y** - has the client subscribed a term deposit? (binary: 'yes', 'no')



Dataset Inspection

AI

```
df.head()
```

| | age | job | marital | education | default | balance | housing | loan | contact | day | month | duration | campaign | pdays | previous | poutcome | y |
|---|-----|--------------|---------|-----------|---------|---------|---------|------|---------|-----|-------|----------|----------|-------|----------|----------|----|
| 0 | 58 | management | married | tertiary | no | 2143 | yes | no | unknown | 5 | may | 261 | 1 | -1 | 0 | unknown | no |
| 1 | 44 | technician | single | secondary | no | 29 | yes | no | unknown | 5 | may | 151 | 1 | -1 | 0 | unknown | no |
| 2 | 33 | entrepreneur | married | secondary | no | 2 | yes | yes | unknown | 5 | may | 76 | 1 | -1 | 0 | unknown | no |
| 3 | 47 | blue-collar | married | unknown | no | 1506 | yes | no | unknown | 5 | may | 92 | 1 | -1 | 0 | unknown | no |
| 4 | 33 | unknown | single | unknown | no | 1 | no | no | unknown | 5 | may | 198 | 1 | -1 | 0 | unknown | no |

```
df.tail()
```

| | age | job | marital | education | default | balance | housing | loan | contact | day | month | duration | campaign | pdays | previous | poutcome | y |
|-------|-----|--------------|----------|-----------|---------|---------|---------|------|-----------|-----|-------|----------|----------|-------|----------|----------|-----|
| 45206 | 51 | technician | married | tertiary | no | 825 | no | no | cellular | 17 | nov | 977 | 3 | -1 | 0 | unknown | yes |
| 45207 | 71 | retired | divorced | primary | no | 1729 | no | no | cellular | 17 | nov | 456 | 2 | -1 | 0 | unknown | yes |
| 45208 | 72 | retired | married | secondary | no | 5715 | no | no | cellular | 17 | nov | 1127 | 5 | 184 | 3 | success | yes |
| 45209 | 57 | blue-collar | married | secondary | no | 668 | no | no | telephone | 17 | nov | 508 | 4 | -1 | 0 | unknown | no |
| 45210 | 37 | entrepreneur | married | secondary | no | 2971 | no | no | cellular | 17 | nov | 361 | 2 | 188 | 11 | other | no |



```
# Checking the head of the numerical features
```

```
df[numerical_features].head()
```

| | age | balance | day | duration | campaign | pdays | previous |
|---|-----|---------|-----|----------|----------|-------|----------|
| 0 | 58 | 2143 | 5 | 261 | 1 | -1 | 0 |
| 1 | 44 | 29 | 5 | 151 | 1 | -1 | 0 |
| 2 | 33 | 2 | 5 | 76 | 1 | -1 | 0 |
| 3 | 47 | 1506 | 5 | 92 | 1 | -1 | 0 |
| 4 | 33 | 1 | 5 | 198 | 1 | -1 | 0 |

After doing the basic dataset inspection we spilled the dataset in categorical and numerical variables separately.

```
# Checking the head of the categorical features
```

```
df[categorical_features].head()
```

| | job | marital | education | default | housing | loan | contact | month | poutcome | y |
|---|--------------|---------|-----------|---------|---------|------|---------|-------|----------|----|
| 0 | management | married | tertiary | no | yes | no | unknown | may | unknown | no |
| 1 | technician | single | secondary | no | yes | no | unknown | may | unknown | no |
| 2 | entrepreneur | married | secondary | no | yes | yes | unknown | may | unknown | no |
| 3 | blue-collar | married | unknown | no | yes | no | unknown | may | unknown | no |
| 4 | unknown | single | unknown | no | no | no | unknown | may | unknown | no |



Data Exploration

- ❖ The dataset has 45211 rows and 17 features (columns).
- ❖ 10 categorical features.
- ❖ 7 Numerical features.
- ❖ No null values.
- ❖ No Duplicate values.
- ❖ No Missing Values.



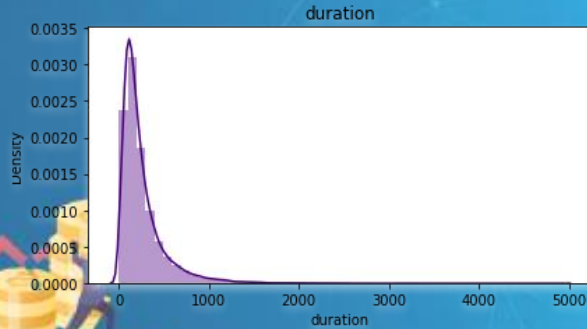
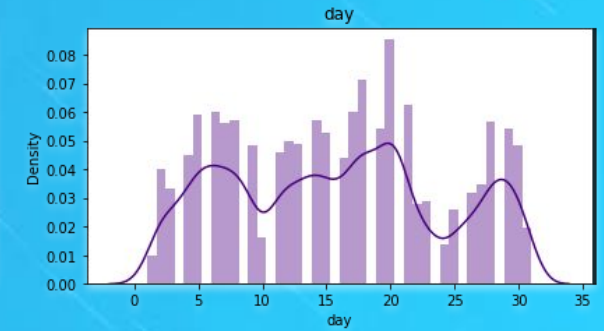
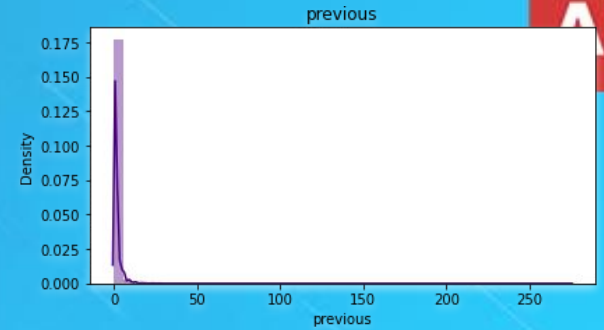

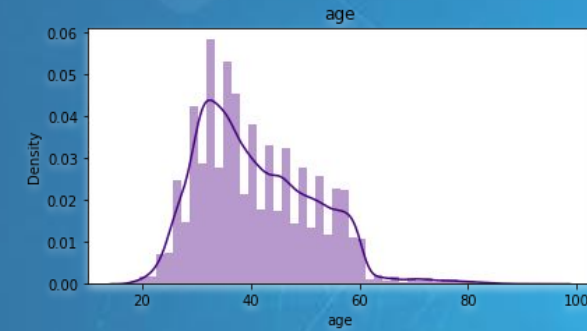


EDA








- ❖ Exploratory data analysis or commonly known as EDA helps to explore data, and possibly formulate hypotheses that might cause new data collection and experiments. EDA build a robust understanding of the data, issues associated with either the info or process. it's a scientific approach to get the story of the data.
- ❖ It focuses more narrowly on checking assumptions required for model fitting and hypothesis testing. It also helps while handling missing values and making transformations of variables as needed.



Univariate analysis of numerical features




Outlier Detection of Numerical Features




Bivariate Analysis of numerical feature with Target Variable

AI


Count on the basis of campaign



Count on the basis of day

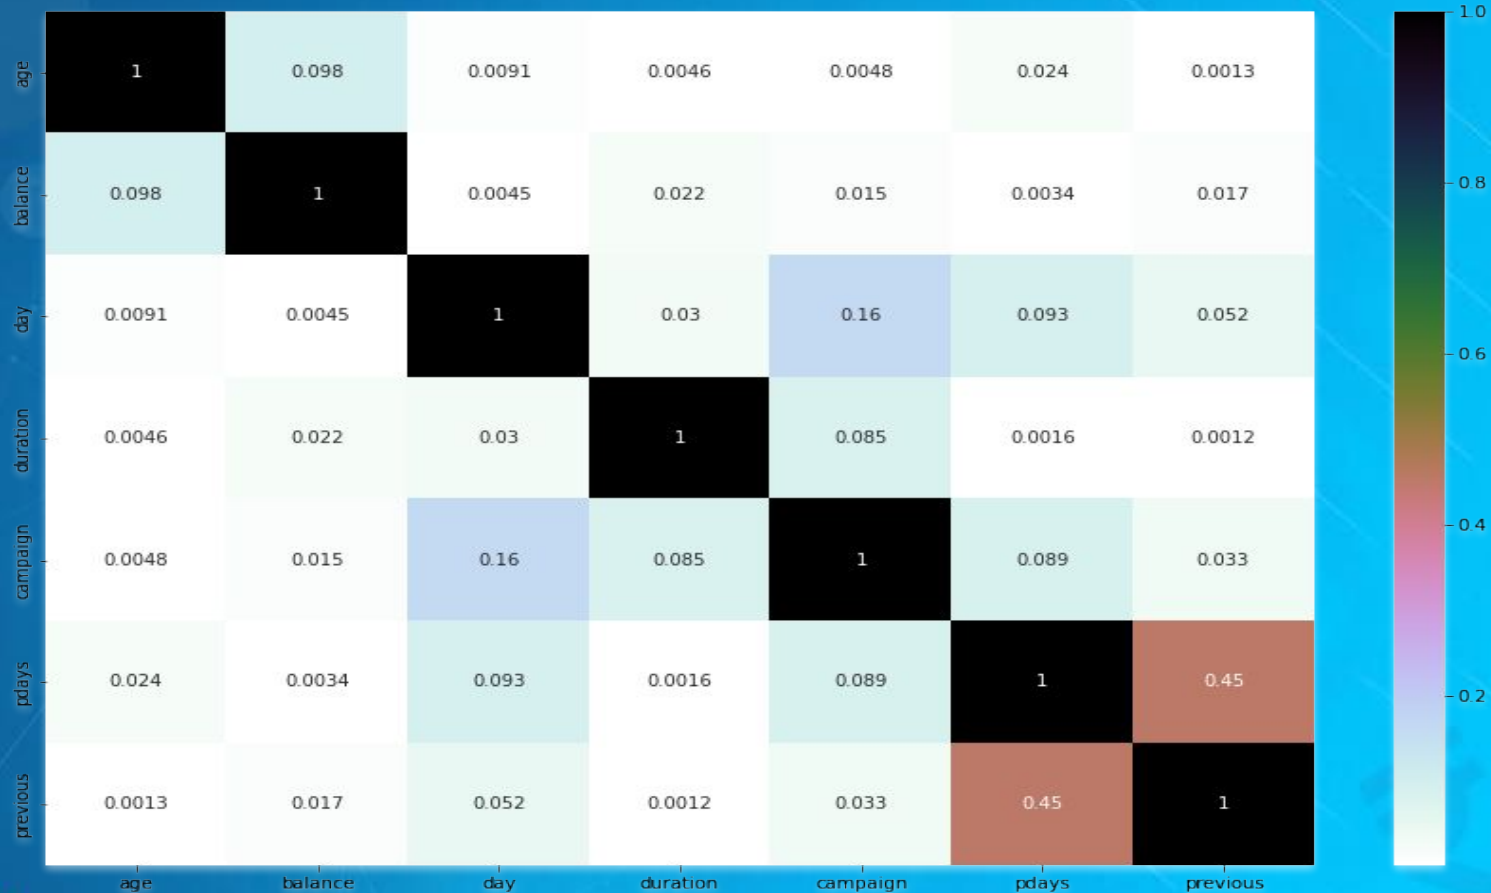


Count on the basis of previous



Checking Multicollinearity

AI



VIF

AI

| | variables | VIF |
|---|-----------|----------|
| 0 | age | 5.004058 |
| 1 | balance | 1.212908 |
| 2 | day | 3.984268 |
| 3 | duration | 1.901309 |
| 4 | campaign | 1.824694 |
| 5 | pdays | 1.454202 |
| 6 | previous | 1.341641 |


- VIF determines the strength of the correlation between the independent variables.
- VIF less than 5 will be included in the model. In some cases VIF of less than 10 is also acceptable.



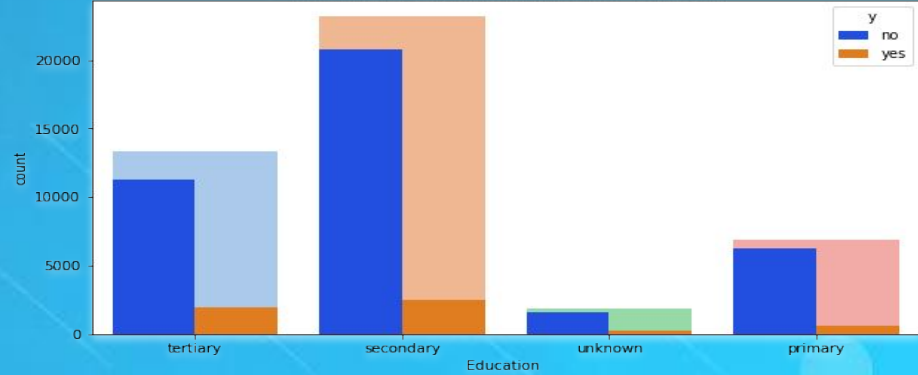
Univariate Analysis of Categorical Features & relation with Target Variable

AI


Count of people with different marital status




Count of yes and no for different education level




Count of different type of jobs



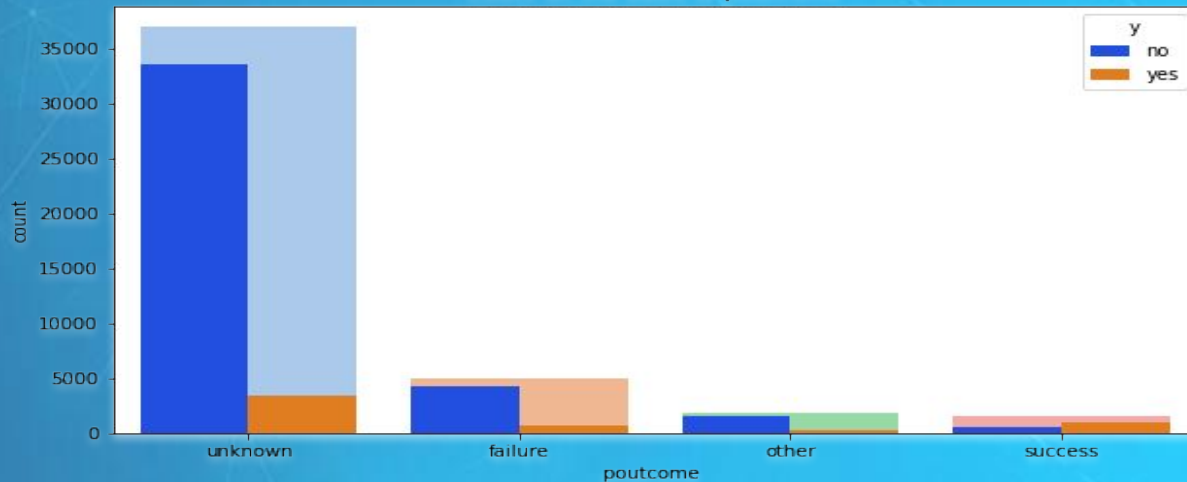
Count of yes and no for different housing




Count of yes and no for different default




Count on the basis of poutcome




Count of yes and no for loan



Count of different types of contact type



Count on the basis of month



Feature Engineering

All machine learning algorithms use some input data to create outputs. Algorithms require features with some specific characteristics to work properly. Here, the need for feature engineering arises. Feature engineering mainly have two goals:

- ❖ Preparing the proper input dataset, compatible with the machine learning algorithm requirements.
- ❖ Improving the performance of machine learning models.

We'll try adding and removing some features in this section in order to make a perfect data matrix we can pass to a machine learning model. We will try to interpret categorical features as numeric to be passed to the ML models.



So, This is the code which we have applied for feature engineering

```
# Getting the dummies of all the categorical features

cat_columns = ['job', 'marital', 'education', 'contact', 'month', 'poutcome']
for col in cat_columns:
    df = pd.concat([df.drop(col, axis=1), pd.get_dummies(df[col],
                                                         prefix=col, prefix_sep='_', drop_first=True,
                                                         dummy_na=False)], axis=1)

# Converting the boolean features in binary

bool_columns = ['default', 'housing', 'loan', 'y']
for col in bool_columns:
    df[col+'_new'] = df[col].apply(lambda x : 1 if x == 'yes' else 0)
    df.drop(col, axis=1, inplace=True)
```

```
# Checking the shape of dataset after all the transformations.
```

```
df.shape
```

```
(45211, 43)
```



Now, Here is what our dataset looks like after all the transformations.

```
df.head()
```


| | age | balance | day | duration | campaign | pdays | previous | job_blue-collar | job_entrepreneur | job_housemaid | ... | month_nov | month_oct | month_sep | poutcome_other |
|---|-----|---------|-----|----------|----------|-------|----------|-----------------|------------------|---------------|-----|-----------|-----------|-----------|----------------|
| 0 | 58 | 2143 | 5 | 261 | 1 | -1 | 0 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 |
| 1 | 44 | 29 | 5 | 151 | 1 | -1 | 0 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 |
| 2 | 33 | 2 | 5 | 76 | 1 | -1 | 0 | 0 | 1 | 0 | ... | 0 | 0 | 0 | 0 |
| 3 | 47 | 1506 | 5 | 92 | 1 | -1 | 0 | 1 | 0 | 0 | ... | 0 | 0 | 0 | 0 |
| 4 | 33 | 1 | 5 | 198 | 1 | -1 | 0 | 0 | 0 | 0 | ... | 0 | 0 | 0 | 0 |

5 rows × 43 columns



Balance of our Target Variable

Proportion of Subscribed & Not Subscribed term Deposit




- As we can see that data is highly imbalanced.
- Majority of the data points belong to “Not - Subscribed” class.

Ratio of “Not - Subscribed” class to “Subscribed” class is 8:1

```
Subscribe
0      39922
1       5289
Name: Subscribe, dtype: int64
```

Handling the imbalance in the dataset using SMOTE



SMOTE (Synthetic Minority Oversampling Technique) works by randomly picking a point from the minority class and computing the k-nearest neighbors for this point. The synthetic points are added between the chosen point and its neighbors.

```
# checking the length of our train set before and after handling imbalance.
```

```
print ('original dataset length',len(X))  
print ('Resample dataset length',len(X_train_sm))
```

```
original dataset length 45211
```

```
Resample dataset length 63884
```



Data Preparation

- ❖ Now that the Dataset is cleaned and we have added all the necessary features along with some conversions of categorical features via.,
 - Label Encoding
 - One Hot Encoding (Dummy Encoding)
- ❖ Then, We used MinMaxscaler for transforming data
- ❖ So, now we have split the data into training and testing sets.
 - Train Test Split (Test size = “0.2” Random state = “0”)



Performance Metrics

- **ROC** also known as Receiver Operating Characteristics, shows the performance of binary class classifiers across the range of all possible thresholds plotting between true positive rate and 1-false positive rate.
- **AUC** measures the likelihood of two given random points, one from positive and one from negative, the classifier will rank the positive points above negative points. AUC-ROC is popular classification metric that presents the advantage of being independent of false positive or negative points.
- **F1 SCORE** is the harmonic mean between Precision and Recall. Macro F1 score is used to know how our model works in overall dataset.
- **Confusion Matrix** gives the count of true negative, true positive, false positive and false negative data points.



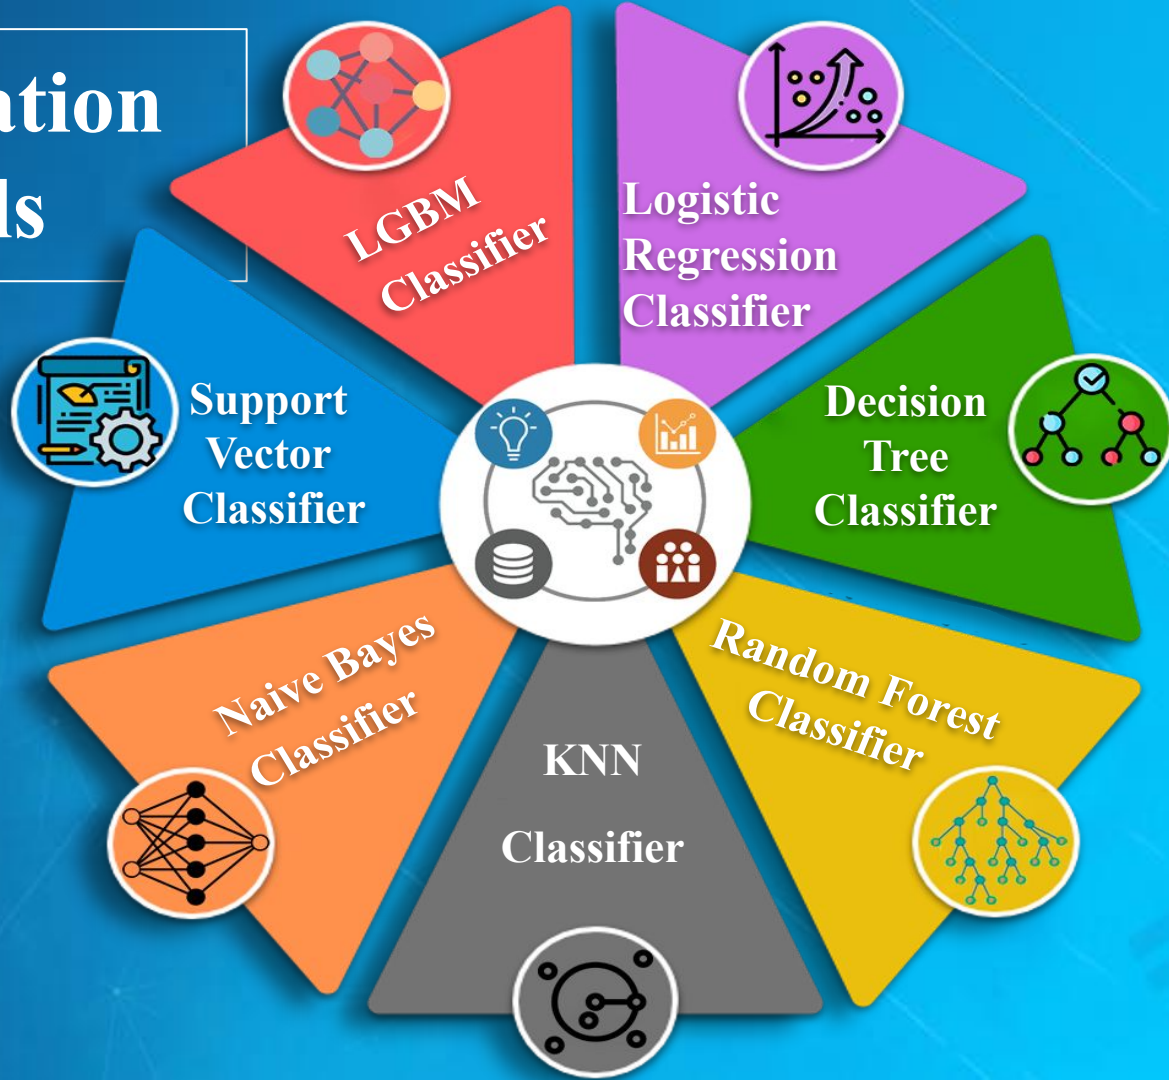
Optimization

- ❖ Hyperparameter tuning consists of finding a set of optimal hyperparameter values for a learning algorithm while applying this optimized algorithm to any data set.
- ❖ GridSearchCV is a technique to search through the best parameter values from the given set of the grid of parameters. It is basically a cross-validation method.



Classification Models

AI




Logistic Regression Classifier

AI

| Precision | Recall | F1 Score | Support |
|-----------|--------|----------|---------|
| 0.8802 | 0.8802 | 0.8802 | 0.8802 |

Confusion Matrix for Logistic Regression



| Accuracy on Train data | Accuracy on Test data |
|------------------------|-----------------------|
| 0.9299 | 0.8802 |




Decision Tree Classifier

AI

| Precision | Recall | F1 Score | Support |
|-----------|--------|----------|---------|
| 0.8824 | 0.8824 | 0.8824 | 0.8824 |

Confusion Matrix for Decision Tree



| Accuracy on Train data | Accuracy on Test data |
|------------------------|-----------------------|
| 0.9214 | 0.8824 |




Random Forest Classifier

AI

| Precision | Recall | F1 Score | Support |
|-----------|--------|----------|---------|
| 0.8821 | 0.8821 | 0.8821 | 0.8821 |

| Accuracy on Train data | Accuracy on Test data |
|------------------------|-----------------------|
| 0.9290 | 0.8821 |

Confusion Matrix for Random Forest




K- Nearest Neighbors Classifier

| Precision | Recall | F1 Score | Support |
|-----------|--------|----------|---------|
| 0.8824 | 0.8824 | 0.8824 | 0.8824 |

| Accuracy on Train data | Accuracy on Test data |
|------------------------|-----------------------|
| 1.0 | 0.8824 |

Confusion Matrix for KNN




Naive Bayes Classifier

AI

| Precision | Recall | F1 Score | Support |
|-----------|--------|----------|---------|
| 0.7863 | 0.7863 | 0.7863 | 0.7863 |

| Accuracy on Train data | Accuracy on Test data |
|------------------------|-----------------------|
| 0.8862 | 0.7863 |

Confusion Matrix for Naive Bayes Classifier




Support Vector Machine Classifier

AI

| Precision | Recall | F1 Score | Support |
|-----------|--------|----------|---------|
| 0.8824 | 0.8824 | 0.8824 | 0.8824 |

| Accuracy on Train data | Accuracy on Test data |
|------------------------|-----------------------|
| 0.6200 | 0.8824 |

Confusion Matrix for Support Vector Classifier




Light Gradient Boost Machine

AI

| Precision | Recall | F1 Score | Support |
|-----------|--------|----------|---------|
| 0.8822 | 0.8822 | 0.8822 | 0.8822 |

| Accuracy on Train data | Accuracy on Test data |
|------------------------|-----------------------|
| 0.9621 | 0.8822 |

Confusion Matrix for LGBM



Evaluation Metrics For All Models

| Model | Test Accuracy | Precision | Recall | F1_score |
|------------------------|---------------|-----------|--------|----------|
| Logistic regression | 0.8802 | 0.8802 | 0.8802 | 0.8802 |
| Decision Tree | 0.8824 | 0.8824 | 0.8824 | 0.8824 |
| Random Forest | 0.8821 | 0.8821 | 0.8821 | 0.8821 |
| K-Nearest Neighbors | 0.8824 | 0.8824 | 0.8824 | 0.8824 |
| Naive Bayes | 0.7863 | 0.7863 | 0.7863 | 0.7863 |
| Support Vector Machine | 0.8824 | 0.8824 | 0.8824 | 0.8824 |
| Light Gradient Boost | 0.8822 | 0.8822 | 0.8822 | 0.8822 |



Conclusions

AI

- ❖ 2nd quarter of the year has the highest number of subscription & Month of May is having the maximum subscriptions.
- ❖ Blue-collar, management and technician showed maximum interest in subscription.
- ❖ Compared to married and single, Divorced people have less interest in term deposit.
- ❖ People with secondary education followed by tertiary education were subscribed to term deposit.
- ❖ Generally people who don't have credit in default are interested in deposit. Majority of the people have home loan but only few of them opted for term deposit.
- ❖ Cellular communication is seen more effective in comparison to other communication types.
- ❖ The calls with large duration has more tendency for conversion.
- ❖ Majority of people were not contacted previously before this campaign.
- ❖ We can choose **KNN** or **Decision Tree** to predict Effectiveness as both of them are showing same accuracy of 88% & F1- Score of (**0.8824**).



Thank you

